



# Persistent Reference: Backend Skill Upgrade (10-Day Plan)

This document is the **single source of truth** for this conversation and future messages. It exists so context is never lost as the chat grows.

---

## User Profile (Technical Context)

- **Primary focus:** Backend Engineering (Node.js)
  - **Secondary:** Backend + DevOps (light, practical)
  - **Not focusing on:** Frontend-heavy work, ML/AI, blogs/docs
  - **Daily time available:** 4–5 hours
  - **Market context:** Bangladesh job market → backend + deployment + infra basics > ML hype
- 

## Current Skill Baseline

### Comfortable With

- Node.js + Express
- REST APIs
- CRUD applications
- JWT authentication
- Role-based access control
- File uploads (Cloudinary)
- MySQL & MongoDB (conceptually + practically)

### Partial / Rusty / Theoretical

- Redis (basic concepts, little real usage)
- WebSockets (learned before, rusty)
- Docker (dockerized once, forgot details)
- CI/CD (theory, no ownership-level hands-on)

### Missing Experience

- Running a backend system in production
  - Async job systems with workers
  - Owning deployment on a VM
  - Connecting infra pieces into one system
-

## Core Project (Locked)

### Project Type

Production-Style Backend API Platform

### Base

- Use **existing hackathon repository** as a foundation
- Reframe from "hackathon submission" → **owned backend service**

### Domain

Async File Processing & Download API

### Core Idea

HTTP is for coordination, not work.

Long-running tasks are handled asynchronously using workers.

---

## Tech Stack (Final)

- Runtime: Node.js
- Framework: Express / Hono (from repo)
- Queue: Redis + BullMQ
- Workers: Separate Node.js worker process
- Storage: MinIO (S3-compatible)
- Caching & Rate Limiting: Redis
- Containerization: Docker + Docker Compose
- CI/CD: GitHub Actions (light, real)
- Deployment: Existing VM (Docker-based)
- Observability: Minimal (metrics > tracing)

## Explicit Constraints (Do NOT Drift)

- ~~No React / frontend UI~~
- ~~No deep OpenTelemetry / Jaeger rabbit holes~~
- ~~No SRE-level infra (read replicas, partitioning, etc.)~~
- ~~No fake scaling claims~~

## End Goal (Day 10 Outcome)

By Day 10, the system should:

- Respond instantly to API requests (<200ms)
- Process long-running jobs via background workers
- Persist job status reliably
- Store results in MinIO and serve via presigned URLs
- Be fully dockerized
- Be deployed on a real VM with a public endpoint
- Have CI checks running on every push

And the user should be able to **explain every design decision calmly in an interview.**

---



## 10-Day Roadmap (High-Level)

- **Day 1:** Remove long HTTP requests → async jobs
  - **Day 2:** Job lifecycle, retries, persistence
  - **Day 3:** MinIO integration (proper storage usage)
  - **Day 4:** Dockerization (API + Worker + Redis + MinIO)
  - **Day 5:** Deployment on VM (ownership moment)
  - **Day 6:** Redis caching + rate limiting
  - **Day 7:** Minimal monitoring (metrics only)
  - **Day 8:** Small-scale load testing & tuning
  - **Day 9:** CI/CD refinement
  - **Day 10:** Polish, README, interview readiness
- 

## Guiding Principle (Always Valid)

Build ONE system. Understand it deeply. Own it end-to-end.

This document should be treated as **authoritative context** for all future guidance in this conversation.